

MSBA 265 - Foundational Module 1: Exploratory Data Analysis & Visual Outlier Audits

Student: Fangqi Hu

Instructor: Shyla Solis

Course: MSBA 265 - Business Analytics Topics

Dataset: French Motor Third Party Liability Claims
(OpenML ID: 41214)

Business Problem Overview

In property and casualty insurance, policy pricing depends on understanding risk exposure, driver age, and population density. Before building predictive loss models or business strategy dashboards, we must inspect raw feature data types, construct a domain-specific Business Data Dictionary, plot feature distributions, evaluate univariate skewness, and visually diagnose extreme outliers.

```
In [17]: import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Set visual styling for clean presentation graphics
plt.style.use('seaborn-v0_8-whitegrid')
plt.rcParams['font.sans-serif'] = 'DejaVu Sans'
plt.rcParams['font.size'] = 10
pd.set_option('display.max_columns', 20)

# Load dataset
data_path = os.path.join("../", "data", "raw_business_data.csv")

if not os.path.exists(data_path):
    data_path = os.path.join("data", "raw_business_data.csv")

df = pd.read_csv(data_path)

print(
    f"[+] Loaded Dataset Shape: "
    f"{df.shape[0]:,} rows x {df.shape[1]} columns"
)

df.head(10)
```

[+] Loaded Dataset Shape: 678,013 rows x 12 columns

Out[17]:

	IDpol	ClaimNb	Exposure	Area	VehPower	VehAge	DrivAge	BonusMalus	VehBrand
0	1.0	1	0.10	'D'	5	0	55	50	'B12'
1	3.0	1	0.77	'D'	5	0	55	50	'B12'
2	5.0	1	0.75	'B'	6	2	52	50	'B12'
3	10.0	1	0.09	'B'	7	0	46	50	'B12'
4	11.0	1	0.84	'B'	7	0	46	50	'B12'
5	13.0	1	0.52	'E'	6	2	38	50	'B12'
6	15.0	1	0.45	'E'	6	2	38	50	'B12'
7	17.0	1	0.27	'C'	7	0	33	68	'B12'
8	18.0	1	0.71	'C'	7	0	33	68	'B12'
9	21.0	1	0.15	'B'	7	0	41	50	'B12'



In [18]: `import sys`
`print(sys.executable)`

c:\Users\87107\msba265_module1\venv\Scripts\python.exe

In [19]: `import sys`
`print(sys.executable)`

c:\Users\87107\msba265_module1\venv\Scripts\python.exe

In [20]: `import os`
`import pandas as pd`
`import numpy as np`
`import matplotlib.pyplot as plt`
`import seaborn as sns`

`plt.style.use('seaborn-v0_8-whitegrid')`
`plt.rcParams['font.sans-serif'] = 'DejaVu Sans'`
`plt.rcParams['font.size'] = 10`
`pd.set_option('display.max_columns', 20)`

`data_path = os.path.join("../", "data", "raw_business_data.csv")`

`if not os.path.exists(data_path):`
 `data_path = os.path.join("data", "raw_business_data.csv")`

`df = pd.read_csv(data_path)`

`print(`
 `f"[+] Loaded Dataset Shape: "`
 `f"{df.shape[0]:,} rows x {df.shape[1]} columns"`
`)`

```
df.head(10)
```

[+] Loaded Dataset Shape: 678,013 rows x 12 columns

```
Out[20]:
```

	IDpol	ClaimNb	Exposure	Area	VehPower	VehAge	DrivAge	BonusMalus	VehBrand
0	1.0	1	0.10	'D'	5	0	55	50	'B12'
1	3.0	1	0.77	'D'	5	0	55	50	'B12'
2	5.0	1	0.75	'B'	6	2	52	50	'B12'
3	10.0	1	0.09	'B'	7	0	46	50	'B12'
4	11.0	1	0.84	'B'	7	0	46	50	'B12'
5	13.0	1	0.52	'E'	6	2	38	50	'B12'
6	15.0	1	0.45	'E'	6	2	38	50	'B12'
7	17.0	1	0.27	'C'	7	0	33	68	'B12'
8	18.0	1	0.71	'C'	7	0	33	68	'B12'
9	21.0	1	0.15	'B'	7	0	41	50	'B12'



```
In [21]: vehage_data = df[["VehAge"]]
display(vehage_data.head(10))
```

	VehAge
0	0
1	0
2	2
3	0
4	0
5	2
6	2
7	0
8	0
9	0

```
In [22]: print("=====")
print("1. STRUCTURAL DATA TYPES & NULLS")
print("=====")
df.info()

print("\n=====")
```

```
print("2. SUMMARY STATISTICS (NUMERICAL)")
print("=====")
df.describe().T
```

```
=====
1. STRUCTURAL DATA TYPES & NULLS
=====
<class 'pandas.DataFrame'>
RangeIndex: 678013 entries, 0 to 678012
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   IDpol            678013 non-null float64
1   ClaimNb          678013 non-null int64
2   Exposure         678013 non-null float64
3   Area             678013 non-null str
4   VehPower         678013 non-null int64
5   VehAge           678013 non-null int64
6   DrivAge          678013 non-null int64
7   BonusMalus       678013 non-null int64
8   VehBrand         678013 non-null str
9   VehGas           678013 non-null str
10  Density          678013 non-null int64
11  Region           678013 non-null str
dtypes: float64(2), int64(6), str(4)
memory usage: 62.1 MB
```

```
=====
2. SUMMARY STATISTICS (NUMERICAL)
=====
```

Out[22]:

		count	mean	std	min	25%	50%	
	IDpol	678013.0	2.621857e+06	1.641783e+06	1.000000	1157951.00	2272152.00	404
	ClaimNb	678013.0	5.324677e-02	2.401173e-01	0.000000	0.00	0.00	
	Exposure	678013.0	5.287501e-01	3.644415e-01	0.002732	0.18	0.49	
	VehPower	678013.0	6.454631e+00	2.050906e+00	4.000000	5.00	6.00	
	VehAge	678013.0	7.044265e+00	5.666232e+00	0.000000	2.00	6.00	
	DrivAge	678013.0	4.549912e+01	1.413744e+01	18.000000	34.00	44.00	
	BonusMalus	678013.0	5.976150e+01	1.563666e+01	50.000000	50.00	50.00	
	Density	678013.0	1.792422e+03	3.958647e+03	1.000000	92.00	393.00	

In [23]: df["VehAge"].describe()


```
Out[23]: count    678013.000000
         mean       7.044265
         std        5.666232
         min         0.000000
         25%        2.000000
         50%        6.000000
         75%       11.000000
         max       100.000000
         Name: VehAge, dtype: float64
```

```
In [24]: import matplotlib.pyplot as plt
import seaborn as sns

sns.histplot(df["VehAge"], bins=25, kde=True)

plt.title("Distribution of Vehicle Age")
plt.xlabel("Vehicle Age (Years)")
plt.ylabel("Count")

plt.show()
```



```
In [25]: business_definitions = {
         "IDpol": (
             "Policy ID",
             "Key ID",
             "Unique policy identification key assigned to contract."
         ),
         "ClaimNb": (
```

```

        "Claim Count",
        "Count",
        "Total insurance claims filed during policy coverage period."
    ),
    "Exposure": (
        "Risk Exposure",
        "Years (0.0-1.0)",
        "Fraction of year policy was active (e.g., 0.50 = 6 months coverage).",
    ),
    "Area": (
        "Area Risk Code",
        "Ordinal Category",
        "Geographic risk code of policyholder address (A = rural, F = urban).",
    ),
    "VehPower": (
        "Engine Power",
        "Tax HP",
        "Fiscal power rating category of vehicle engine."
    ),
    "VehAge": (
        "Vehicle Age",
        "Years",
        "Age of the insured vehicle in years."
    ),
    "DrivAge": (
        "Driver Age",
        "Years",
        "Age of the primary insured driver in years."
    ),
    "BonusMalus": (
        "Claim History Index",
        "Index Score (50-350)",
        "Regulatory driver risk score (<100: safe discount; >100: claim surcharge).",
    ),
    "VehBrand": (
        "Brand Class",
        "Category Code",
        "Anonymized vehicle manufacturer / brand classification code."
    ),
    "VehGas": (
        "Fuel Type",
        "Binary Class",
        "Engine fuel type used (Diesel vs. Regular / Gasoline).",
    ),
    "Density": (
        "Population Density",
        "People / km2",
        "Population density of policyholder's municipality."
    ),
    "Region": (
        "Administrative Region",
        "Region Code",
        "French administrative region code corresponding to residency location."
    )
}

```

```

dict_list = []

for col in df.columns:
    label, unit, definition = business_definitions.get(
        col,
        (col, "N/A", "No business definition provided.")
    )

    dict_list.append({
        "Feature Name": col,
        "Business Label": label,
        "Data Type": str(df[col].dtype),
        "Unit": unit,
        "Null Count": df[col].isnull().sum(),
        "Unique Values": df[col].nunique(),
        "Business Definition": definition
    })

dict_df = pd.DataFrame(dict_list)

display(dict_df)

reports_dir = os.path.join("../", "reports")
os.makedirs(reports_dir, exist_ok=True)

dict_path = os.path.join(reports_dir, "data_dictionary.csv")
dict_df.to_csv(dict_path, index=False)

print(f"\n[+] Successfully exported Business Data Dictionary artifact to: {dict_pat

```

	Feature Name	Business Label	Data Type	Unit	Null Count	Unique Values	Business Definition
0	IDpol	Policy ID	float64	Key ID	0	678013	Unique policy identification key assigned to c...
1	ClaimNb	Claim Count	int64	Count	0	11	Total insurance claims filed during policy cov...
2	Exposure	Risk Exposure	float64	Years (0.0-1.0)	0	181	Fraction of year policy was active (e.g., 0.50...
3	Area	Area Risk Code	str	Ordinal Category	0	6	Geographic risk code of policyholder address (...)
4	VehPower	Engine Power	int64	Tax HP	0	12	Fiscal power rating category of vehicle engine.
5	VehAge	Vehicle Age	int64	Years	0	78	Age of the insured vehicle in years.
6	DrivAge	Driver Age	int64	Years	0	83	Age of the primary insured driver in years.
7	BonusMalus	Claim History Index	int64	Index Score (50-350)	0	115	Regulatory driver risk score (<100: safe disco...
8	VehBrand	Brand Class	str	Category Code	0	11	Anonymized vehicle manufacturer / brand classi...
9	VehGas	Fuel Type	str	Binary Class	0	2	Engine fuel type used (Diesel vs. Regular / Ga...
10	Density	Population Density	int64	People / km2	0	1607	Population density of policyholder's municipal...
11	Region	Administrative Region	str	Region Code	0	22	French administrative region code correspondin...

[+] Successfully exported Business Data Dictionary artifact to: ..\reports\data_dictionary.csv

Business Interpretation: Exposure and Spatial Features

Exposure: Exposure represents the fraction of time that a policyholder was actually insured and at risk. Therefore, raw ClaimNb should not be interpreted without considering exposure time. A policyholder insured for one year has more opportunity to file a claim than a policyholder insured for only several months. Claim frequency should therefore be evaluated relative to Exposure, or Exposure can later be used as an offset in an appropriate count model.

Density vs. Area: Density and Area both contain information about geographic environment and urbanization. Density is a continuous measure of population concentration, while Area is a categorical or ordinal geographic risk code. Including both without justification may introduce redundant spatial information and make model coefficients harder to interpret. Therefore, feature selection should be based on business meaning and correlation diagnostics rather than automatically including both variables.

In [26]: *# My assigned feature: VehAge*

```
vehage_data = df[["VehAge"]]  
  
display(vehage_data.head(10))
```

VehAge	
0	0
1	0
2	2
3	0
4	0
5	2
6	2
7	0
8	0
9	0

```
In [27]: print("=====")  
print("UNIVARIATE SKEWNESS DIAGNOSTICS")  
print("=====")  
  
numeric_df = df.select_dtypes(include=[np.number])  
  
skewness_series = numeric_df.skew().sort_values(ascending=False)  
  
skew_summary = pd.DataFrame({  
    "Feature": skewness_series.index,  
    "Skewness Coefficient": skewness_series.values,
```

```

    "Diagnosis": [
        "Extremely Right-Skewed" if s > 3.0
        else "Moderately Skewed" if abs(s) > 1.0
        else "Unskewed / Symmetric"
        for s in skewness_series.values
    ],
    "Action Required": [
        "Log / Box-Cox Transform" if abs(s) > 1.0
        else "Standard Scaler Sufficient"
        for s in skewness_series.values
    ]
})

display(skew_summary)

```

```

=====
UNIVARIATE SKEWNESS DIAGNOSTICS
=====

```

	Feature	Skewness Coefficient	Diagnosis	Action Required
0	ClaimNb	5.599613	Extremely Right-Skewed	Log / Box-Cox Transform
1	Density	4.651421	Extremely Right-Skewed	Log / Box-Cox Transform
2	BonusMalus	1.728934	Moderately Skewed	Log / Box-Cox Transform
3	VehPower	1.171344	Moderately Skewed	Log / Box-Cox Transform
4	VehAge	1.147989	Moderately Skewed	Log / Box-Cox Transform
5	DrivAge	0.435759	Unskewed / Symmetric	Standard Scaler Sufficient
6	IDpol	0.237890	Unskewed / Symmetric	Standard Scaler Sufficient
7	Exposure	0.085318	Unskewed / Symmetric	Standard Scaler Sufficient

ClaimNb Transformation Decision

Although ClaimNb is strongly right-skewed and an automated diagnostic may recommend a logarithmic transformation, I would override that recommendation.

ClaimNb is a discrete count variable, and many policyholders have exactly zero claims. A direct logarithmic transformation is mathematically undefined for zero because $\log(0)$ approaches negative infinity. More importantly, ClaimNb represents count events rather than a continuous positive volume.

Therefore, ClaimNb should remain in its original count form and should later be modeled using an appropriate count or rate approach rather than being blindly log-transformed.

```

In [28]: # My assigned feature: VehAge

vehage_data = df[["VehAge"]]

```

```
display(vehage_data.head(10))
```

VehAge	
0	0
1	0
2	2
3	0
4	0
5	2
6	2
7	0
8	0
9	0

```
In [29]: print("=====")
print("PEARSON MULTICOLLINEARITY AUDIT")
print("=====")

corr_matrix = numeric_df.corr().abs()

upper_tri = corr_matrix.where(
    np.triu(np.ones(corr_matrix.shape), k=1).astype(bool)
)

high_corr_pairs = []

for col in upper_tri.columns:
    high_corr_features = upper_tri.index[
        upper_tri[col] >= 0.85
    ].tolist()

    for feat in high_corr_features:
        val = upper_tri.loc[feat, col]

        high_corr_pairs.append({
            "Feature 1": feat,
            "Feature 2": col,
            "Pearson Correlation (|r|)": round(val, 4),
            "Status": "ACTION REQUIRED: Severe Multicollinearity"
        })

if high_corr_pairs:
    display(pd.DataFrame(high_corr_pairs))
else:
    print(
        "[+] No severe collinear feature pairs "
```

```
)    "(|r| >= 0.85) detected across numerical metrics."
)
```

```
=====
PEARSON MULTICOLLINEARITY AUDIT
=====
```

```
[+] No severe collinear feature pairs (|r| >= 0.85) detected across numerical metrics.
```

In [30]: *# Cell 6b: Render Annotated Pearson Correlation Heatmap*

```
plt.figure(figsize=(10, 8))

# Compute Pearson correlation matrix across numerical variables
corr = numeric_df.corr()

# Generate correlation heatmap
sns.heatmap(
    corr,
    annot=True,
    fmt=".2f",
    cmap="coolwarm",
    vmin=-1.0,
    vmax=1.0,
    linewidths=0.5,
    cbar_kws={"label": "Pearson Correlation Coefficient (r)"}
)

plt.title(
    "Pearson Correlation Heatmap: Numerical Features",
    fontsize=14,
    fontweight="bold",
    pad=12
)

plt.tight_layout()

heatmap_path = os.path.join(
    "..",
    "reports",
    "figures",
    "correlation_heatmap.png"
)

plt.savefig(heatmap_path, dpi=300)
plt.show()

print(
    f"[+] Saved correlation heatmap graphics artifact to: {heatmap_path}"
)
```



[+] Saved correlation heatmap graphics artifact to: `..\reports\figures\correlation_heatmap.png`

In []:

Pearson Correlation Heatmap Interpretation

- **Distribution Shape:** The correlation matrix is dominated by weak linear relationships across the numerical features. No feature pair shows severe multicollinearity.
- **Key Metrics / Whiskers:** The strongest relationship is between DrivAge and BonusMalus, with a moderate negative correlation of approximately $r = -0.48$. No numerical feature pair reaches the severe multicollinearity threshold of $|r| \geq 0.85$.
- **Recommended Analytics Action:** Retain the numerical features because none requires removal based solely on Pearson multicollinearity. However, potential business redundancy between Area and Density should still be evaluated using domain knowledge.

```
In [31]: fig, axes = plt.subplots(2, 2, figsize=(14, 10))

# 1. Population Density Box Plot
```

```

sns.boxplot(
    x=df["Density"],
    ax=axes[0, 0],
    color="lightcoral"
)
axes[0, 0].set_title(
    "Box Plot: Population Density (People / km2)",
    fontsize=12,
    fontweight="bold"
)
axes[0, 0].set_xlabel("Density")

# 2. Driver Age Histogram + KDE
sns.histplot(
    df["DrivAge"],
    kde=True,
    ax=axes[0, 1],
    color="skyblue",
    bins=30
)
axes[0, 1].set_title(
    "Histogram & KDE: Driver Age Distribution",
    fontsize=12,
    fontweight="bold"
)
axes[0, 1].set_xlabel("Driver Age (Years)")

# 3. BonusMalus Box Plot
sns.boxplot(
    x=df["BonusMalus"],
    ax=axes[1, 0],
    color="mediumpurple"
)
axes[1, 0].set_title(
    "Box Plot: BonusMalus Risk Index",
    fontsize=12,
    fontweight="bold"
)
axes[1, 0].set_xlabel("BonusMalus Index Score")

# 4. Vehicle Age Histogram + KDE
sns.histplot(
    df["VehAge"],
    kde=True,
    ax=axes[1, 1],
    color="seagreen",
    bins=25
)
axes[1, 1].set_title(
    "Histogram & KDE: Vehicle Age Distribution",
    fontsize=12,
    fontweight="bold"
)
axes[1, 1].set_xlabel("Vehicle Age (Years)")


plt.tight_layout()

```

```
fig_path = os.path.join(
    "..",
    "reports",
    "figures",
    "feature_distributions.png"
)

plt.savefig(fig_path, dpi=300)
plt.show()

print(
    f"[+] Saved distribution graphics artifact to: {fig_path}"
)
```



[+] Saved distribution graphics artifact to: ..\reports\figures\feature_distributions.png

In []:

Visual Distribution & Outlier Interpretation

1. Population Density (Density)

- **Distribution Shape:** Density is extremely right-skewed. Most observations are concentrated at relatively low density levels, while a long right tail extends to approximately 27,000 people/km².

- **Key Metrics / Whiskers:** Q1 $\approx$ 92, median $\approx$ 393, Q3 $\approx$ 1,658, and maximum $\approx$ 27,000. Using the Tukey IQR rule, the upper fence is approximately 4,007, so many observations above this value appear as statistical outliers.
- **Recommended Analytics Action:** These high-density observations may represent valid policyholders living in major urban areas rather than data-entry errors. They should not be removed automatically. A transformation such as $\log(1 + \text{Density})$ may be more appropriate for downstream modeling.

2. Driver Age (DrivAge)

- **Distribution Shape:** Driver age is relatively stable and approximately unimodal, with most observations concentrated around middle-aged drivers.
- **Key Metrics / Whiskers:** The observed range is approximately 18 to 100 years, with Q1 $\approx$ 34, median $\approx$ 44, and Q3 $\approx$ 55.
- **Recommended Analytics Action:** The values remain within plausible human and legal boundaries, so no outlier removal is required. Standard scaling may be applied later if required by the modeling method.

3. BonusMalus Risk Index

- **Distribution Shape:** BonusMalus is right-skewed. Most policyholders are concentrated around lower risk-score values, while a smaller group has substantially higher scores.
- **Key Metrics / Whiskers:** Q1 $\approx$ 50, median $\approx$ 50, Q3 $\approx$ 64, and the observed maximum is approximately 230. High-score observations appear beyond the upper whisker.
- **Recommended Analytics Action:** High BonusMalus values may represent legitimate higher-risk or surcharge drivers rather than errors. They should not be removed automatically. A business indicator such as `Is_Surcharge_Driver = BonusMalus > 100` could be considered for later modeling.

4. Vehicle Age (VehAge)

- **Distribution Shape:** Vehicle age is right-skewed. Most insured vehicles are relatively young, while a smaller number of much older vehicles create a long right tail.
- **Key Metrics / Whiskers:** Q1 $\approx$ 2 years, median $\approx$ 6 years, Q3 $\approx$ 11 years, and the observed maximum is approximately 100 years.
- **Recommended Analytics Action:** Very old vehicles should be investigated using domain knowledge before removal. They may represent rare but valid records. Extreme values should therefore be audited rather than automatically deleted.

```
In [32]: target_feature = "Density"

# Tukey IQR method
Q1 = df[target_feature].quantile(0.25)
Q3 = df[target_feature].quantile(0.75)
IQR = Q3 - Q1
```

```

iqr_lower = Q1 - 1.5 * IQR
iqr_upper = Q3 + 1.5 * IQR

df_iqr_cleaned = df[
    (df[target_feature] >= iqr_lower) &
    (df[target_feature] <= iqr_upper)
]

# Z-score method
z_scores = (
    df[target_feature] - df[target_feature].mean()
) / df[target_feature].std()

df_z_cleaned = df[abs(z_scores) <= 3.0]

# Report
print(f"=== OUTLIER AUDIT REPORT FOR '{target_feature}' ===")
print(f"Original Record Count: {len(df):,} records")

print(
    f"Tukey IQR Range [{iqr_lower:.2f}, {iqr_upper:.2f}]: "
    f"Excludes {len(df) - len(df_iqr_cleaned):,} records"
)

print(
    f"Z-Score Threshold (|Z| <= 3.0): "
    f"Excludes {len(df) - len(df_z_cleaned):,} records"
)

# Before / After visualization
fig, axes = plt.subplots(1, 2, figsize=(14, 5))

sns.boxplot(
    x=df[target_feature],
    ax=axes[0],
    color="salmon"
)
axes[0].set_title(
    f"BEFORE Filtering: Raw {target_feature}",
    fontsize=12,
    fontweight="bold"
)
axes[0].set_xlabel("Population Density (People / km2)")

sns.boxplot(
    x=df_iqr_cleaned[target_feature],
    ax=axes[1],
    color="lightgreen"
)
axes[1].set_title(
    f"AFTER Filtering: IQR Tukey Filtered {target_feature}",
    fontsize=12,
    fontweight="bold"
)
axes[1].set_xlabel("Population Density (People / km2)")

```

```

plt.tight_layout()

fig_filter_path = os.path.join(
    "..",
    "reports",
    "figures",
    "outlier_filtering_comparison.png"
)

plt.savefig(fig_filter_path, dpi=300)
plt.show()

print(
    f"[+] Saved visual outlier comparison artifact to: "
    f"{fig_filter_path}"
)


```

=== OUTLIER AUDIT REPORT FOR 'Density' ===

Original Record Count: 678,013 records

Tukey IQR Range [-2257.00, 4007.00]: Excludes 77,566 records

Z-Score Threshold ($|Z| \leq 3.0$): Excludes 15,209 records



[+] Saved visual outlier comparison artifact to: ..\reports\figures\outlier_filtering_comparison.png

Outlier Filtering Interpretation

The outlier audit evaluated population density (Density) across 678,013 policy records.

Using the non-parametric Tukey IQR method, the valid range was approximately -2,257 to 4,007 people/km². This method identified 77,566 records (about 11.44%) as statistical outliers.

In contrast, the parametric Z-score method ($|Z| > 3.0$) excluded only 15,209 records (about 2.24%). The large difference occurs because Density is strongly right-skewed. Extreme urban-density observations increase the mean and standard deviation, causing the Z-score threshold to become much less restrictive.

From a business perspective, many high-density observations likely represent legitimate urban policyholders rather than data-entry errors. Therefore, automatically deleting all Tukey

outliers could remove a meaningful portion of the customer population and introduce selection bias.

If the business objective is statewide insurance pricing, it may be better to retain valid urban records and apply a transformation such as $\log(1 + \text{Density})$. If the downstream model is highly sensitive to extreme leverage points, Tukey trimming may be considered.